

Project report - Frequent Pattern Mining

Università Ca' Foscari Venezia - Computer Science MSc

Matteo De Giorgi (872029)

Problem description

The problem that has been addressed is mining of frequent itemsets, where an itemset is the collection of one or more items (e.g. $\{1, 2, 4, 10\}$), which is considered frequent when it appears in more than a given percentage of the input itemsets.

The input is composed by:

- The path to a file containing new-line separated lists items like the following:

```
1 2 3 4 5 6 7 8
8 7 6
6
2 6 7 1
1 6 7 8 10 11
5 6 7 8
1 2 3 4 6 7 8
6 7 8
```

- A number representing the percentage of transactions that need to contain an itemset in order to be considered frequent.

For example, using a fraction percentage of 50% the result over the previously reported example is composed by 11 itemsets as follows:

```
7
7 6
8
8 6
8 7
8 7 6
1
1 6
1 7
1 7 6
6
```

Sequential solution

In order to provide the first sequential solution it has been decided to use an approach based on Frequent Pattern Trees (FP-Trees), which have been proven to be very effective for representing a compressed version of the transaction database (the input file). In an FP-Tree each transaction is represented by a path from the root to a leaf, moreover a header table is used in order to keep track of the first node for a certain item, which is then linked to all the subsequent ones in a list fashion.

First sequential implementation

The first approach to the sequential algorithm has been tried simplifying the canonical algorithm described above avoiding pruning the FP-Tree between recursive calls, and hence without calling the functions `removeItem` and `pruneInfrequent`, but instead, removing only the reference to the infrequent items by the header table.

This approach reduces the algorithmic complexity since removing the prefix item, and the infrequent items only from the headerTable have constant time complexity per removal operation (given its implementation as associative map). Instead, effectively removing them from the tree has a bigger computational cost since it implies starting from the first node pointed by the header table and iterating over every item in the list. Moreover, it is possible that the removal of a node triggers a cascade of join operations between nodes, since the children of a removed item can have the same value as one of its uncles.

The downside of this approach is the huge memory usage, since the size of the FP-Tree stays the same from the beginning until the end. This implies that in the worst case, when all the possible combinations of the items are frequent, we will need to store 2^n copies of the initial FP-Tree, where n is the number of unique items.

Despite the more the list of prefixes is expanded, and hence the deeper the recursion into **FP-growth**, the lesser is likely that the sequence can be considered frequent, hence it is not likely to approach the maximum theoretical upper limit of 2^n copies. It is still an extremely high upper bound and after experimental tests it has proven to make this approach not feasible nor efficient.

Second sequential implementation

The second approach includes the previously missing functions **removeItem** and **pruneInfrequent**. This immediately denoted a heavy reduction of memory usage, which can be explained by the fact that for every chosen prefix item the prefix FP-Tree associated to it now does not contain any infrequent item nor the chosen prefix item.

Since the previous implementation highlighted how important it is to make an efficient use of memory, this implementation required a particular attention to the FP-Tree nodes pointers' ownership. Since it is easy to create reference cycles, which would imply that nodes would not be actually deallocated causing memory leaks.

In order to solve this problem it has been decided that every FP-Tree node would contain strong references (**shared_ptr**) only to his children and keep all the other references (to the parent or the next and previous node with the same value) as weak references (**weak_ptr**).

Moreover, in order to implement the list of nodes with the same value it has been decided to use a doubly linked list. This has the advantage to reduce the complexity of the removal operation of a specific node, since in order to find the previous one in the list it is not necessary to start from the first node in the header table. This may not sound necessary at first, since the removal operations that we need iterate over the exact list of nodes beginning from the header table, instead it has proven to be a very useful choice since removing a node may result in having to merge one of his children with one of its uncles.

Parallel decomposition

In order to parallelize the program it has been used OpenMP which provides a portable directive-based API in order to offer to the programmer with a higher level of abstraction than pthreads.

First parallel implementation

The first approach to the code parallelization has been to use the OpenMP for-loop directive in the main loop of the **FP-growth** algorithm implementation. This implies splitting the computation of the prefix trees over multiple threads which resulted in a good initial performance improvement. For this implementation it has been necessary to multi-thread only the main loop in **FP-growth** and not all the nested ones, which otherwise, would generate an exponential amount of threads, hence resulting in a heavy degradation of performance because of their overhead (including the risk of reaching the maximum amount of threads in the system aborting the program execution).

Moreover, thanks to the coverage analysis, it has been determined that a big part of the execution time it is spent in the function that deletes infrequent items. This function first deletes the item from the header

table, and then iterates over the list of items with the same value and removes them from the tree. This last step has been proven quite time-consuming, especially over big datasets with many infrequent items.

Unfortunately, it has not been possible to parallelize this cycle iterating over all the elements using the OpenMP for-loop directive, which requires to have a for loop expressed in the “Canonical Loop Form” and hence using a random-access iterator. This because the implementation uses a manually implemented list, in fact every FP-Tree node, besides references to his parent and children, contains also references to his predecessor and successor in this list. Therefore, only a bidirectional iterator could be created, hence the cycle has been parallelized using the OpenMP explicit tasks subdivision.

Handling OpenMP tasks manually is slower than using a for-loop, since the latest subdivides the load equally among the items (assuming a similar workload for each item), and the second assigns each iteration to a task which then needs to be scheduled to a thread, hence having a higher overhead.

This last change required to ensure that nodes’ frequencies are not updated concurrently by multiple threads, hence it has been added OpenMP atomic sections to the arithmetic operations on frequencies. Moreover, since now multiple deletions can happen concurrently it is necessary to ensure that operations on the underlying STL map data structure holding the header table are thread safe. In order to do that an OpenMP critical section has been added ensuring that only a single thread at a time can update the header table.

Second parallel implementation

The first implementation has soon begun to present his main issues during experimental testing. Despite it is effective in some datasets, it can easily result in an uneven load distribution among the threads, for instance, where there is one element that appears in every itemset and has multiple children. In this case only 1 thread would be generated, since the for-loop was prevented to spawn nested threads, hence having the worst balancing possible.

In order to solve this issue, the way the for-loop in **FP-growth** was parallelized has been rethought in order to have a better distribution of the load also in nested recursive calls. That has been achieved using the OpenMP directive **taskloop**, which is able to split the for loop in chunks and create a task for each of them, which will be then assigned to the first available thread, where the amount of maximum concurrent threads is limited via command line parameter.

Thanks to this new approach there is a better load balancement since every recursive step will be able to leverage on multi-threading, the trade-off to accept is the increased overhead due to tasks scheduling. Multiple tests have been performed in order to have a comparison between the approach based on the OpenMP **for** and **taskloop** directives. Which show that they produce comparable results in datasets that were not problematic for the first parallel implementation, and provide a very good improvement in the others.

Another improvement made in this implementation is the parallelization of the function to recompute the support once a new FP-Tree projection needs to be computed (hence everytime a new prefix is selected in **FP-growth**). In order to do that OpenMP manual tasks have been used again, hence when iterating over the list of nodes where the new support is computed from, a new task is created for each of them.

Third parallel implementation

After many experimental testing, it has been realised that some introduced parallel implementations have actually decreased performances on certain datasets. For example the manual tasks creation during the support re-computation has proven to have very little if not a negative effect on performances and hence it has been removed.

Moreover, the parallel implementation of **deleteItem** has proved to be quite debatable, since it is very helpful when a lot of deletions takes place, for instance using high supports or datasets with many infrequent items. At the same time it causes overhead generating more tasks when the main effort is not

deleting items, but exploring all the possible prefixes and hence in the **FP-growth** algorithm. Hence, it is useful only in case the load needs to be balanced and there is a lack of executable tasks.

The parallelization of the items deletion function has been object to multiple tests, which resulted in its conditional inclusion, hence offloading the best decision to the user since it is not possible to know how the parallel implementation will behave beforehand.

Another important change was the parallelization of the initial computation of the FP-Tree, which has been done using manual **task** generation. Previously this was the only phase of the program that was not taking advantage of multithreading. Therefore, now every ordered sequence read from the input dataset is processed by a different task and hence potentially by different threads.

Critical sections added as part of the first parallel implementation has been changed in order to general purpose locks. This has had the effect of reducing useless waits since now the lock is specific to an object instance and not just related to the code section like before, hence, allowing multiple threads to be in the same critical section if it is related to different objects.

This has proved particularly useful during the deletion of items, since operations on the header table can be performed concurrently if relative to two different tables. Moreover, it affects also the computation of the initial FP-Tree, which allows the modification of multiple nodes at the same time and restricts only concurrent modifications to the same node.

Results

After the third parallel implementation, many tests have been performed in order to verify the assumptions that motivated the parallelization choices. In particular, it has been tested what are the limits of the implemented multithreading approach checking execution time and how the load is subdivided among the various threads using various datasets.

One of the main goal of the parallel implementation is to ensure that the load is as balanced as possible among the various threads. This has been tested making sure that the tasks queue is always populated and hence that the CPU utilization stays high along the whole execution of the program.

Experimental results has shown that there is a good subdivision of the load, and hence for most of the datasets the algorithm is able to keep the maximum amount of threads busy. There are some exceptions for datasets that include a lot of infrequent elements, these cases can be improved using the parallelized deletion (using the **-e** option). The support fraction chosen by the user has a heavy impact on this, in fact it influences the amount of items that requires pruning, hence it is necessary to find a good balance through experiments.

Unfortunately, there are still cases where not enough tasks can be produced in order to keep all the threads busy, this is a consequence of the fact that there are many different item lists that needs to be deleted but only one list can be deleted at a time. The reason behind this is that concurrent deletions of different lists could interfere in too many ways to actually provide a performance improvement. Hence, in the case where many short item lists needs to be pruned from the tree it is going to be difficult to generate enough tasks.

It has been tested how performance increases varying the amount of threads, and, as expected, we can always see a performance improvement increasing the amount of threads. Up to the point where the number of CPU cores is reached, from then on we can see a stabilization that eventually leads to a performance decrease due to the context switch overhead. Below it is possible to see a plot of the execution time (in seconds) increasing the amount of threads. These tests have been performed on an 8 core machine.

As we can see the lowest execution time corresponds to the usage of 8 threads (which match the number of CPU cores).

The final average measured Speed-Up is 1.5, using 8 threads, which is much lower than the expected result. This because the algorithm parallelization contains a lot of constraints to ensure data structures




Figure 1: Varying execution time based on the amount of threads

integrity, since many operations cannot be parallelized completely.

Future improvements

The final parallel version has been analyzed using `perf` in order to read CPU performance counters during various common runs.

It has been noticed that `LLC-load-misses:u`, expressed as a percentage the number of loads that miss in the last level cache (typically the L3 for modern Intel chips), can be quite high (around 30%), especially when using low supports. This means that the usage of cache could be improved, especially in all those cases when chain of pointers has been used, since they do not take advantage of the principles of spatial and temporal locality.

Instead, it would be interesting to investigate different algorithms that would allow making threads cooperate for the cache usage, since now they are all trying to load different sub-trees of the same parent. Moreover, the used data structures are not always optimal to do this, since an FP-Tree is a collection of pointers and the whole tree cannot always fit in memory.